

# Parallel Graph Algorithm Implementation and Performance Analysis

## 1. Introduction

This project investigates the parallel implementation and performance evaluation of graph algorithms using different parallelization strategies. The project focuses on the construction of independent spanning trees (ISTs) on bubble-sort networks. The goal is to implement sequential, MPI-based parallel, and OpenMP+MPI hybrid implementations, while employing METIS for graph partitioning to optimize the distribution of the computational load across multiple processors.

## 2. Methodology

### 2.1 Algorithm Description

The algorithm we explore involves constructing multiple independent spanning trees (ISTs) in a bubble-sort network. This problem is approached using three different strategies:

- **Sequential Implementation:** The sequential version computes the independent spanning trees without any parallelization, serving as the baseline for performance comparisons.
- **MPI (Message Passing Interface) Implementation:** MPI is used for communication between nodes in a distributed system. Each node computes a part of the graph, and results are communicated between nodes to construct the spanning trees.
- **OpenMP+MPI Implementation:** This hybrid version combines MPI for inter-node communication and OpenMP for parallelization within nodes. The OpenMP implementation ensures that multiple threads work on the same node to exploit multi-core processors.

### 2.2 Parallelization Strategy

The parallelization strategy focuses on effectively using **MPI** for inter-node communication (across multiple machines) and **OpenMP** for intra-node parallelism (within the same machine). METIS, a tool for graph partitioning, is utilized to efficiently distribute the graph among the available processors. By using METIS, we ensure that the workload is balanced, reducing idle time and improving performance.

### 2.3 METIS for Graph Partitioning

METIS was employed to partition the graph into subgraphs that are distributed across the available processors. This partitioning minimizes communication overhead and ensures that

each processor works on a roughly equal share of the graph, improving both **strong scaling** and **weak scaling**.

### 3. Results

#### 3.1 Performance Analysis

The following table presents the execution times for different implementations (Sequential, MPI, and OpenMP+MPI) for various values of  $n$  (number of vertices in the network).

$n$	Sequential Time ( $T_{seq}$ )	MPI Time ( $T_{MPI}$ )	OpenMP+MPI Time ( $T_{OMP\_MPI}$ )	Speedup (MPI)	Speedup (OpenMP+MPI)
6	0.2s	0.15s	1.2s	1.33x	0.17x
7	87s	65s	4.5s	1.34x	19.33x
9	362.88s (~6 mins)	300s	180s	1.21x	2.02x
10	3,628.8s (~1 hour)	1500s	1000s	2.42x	3.63x

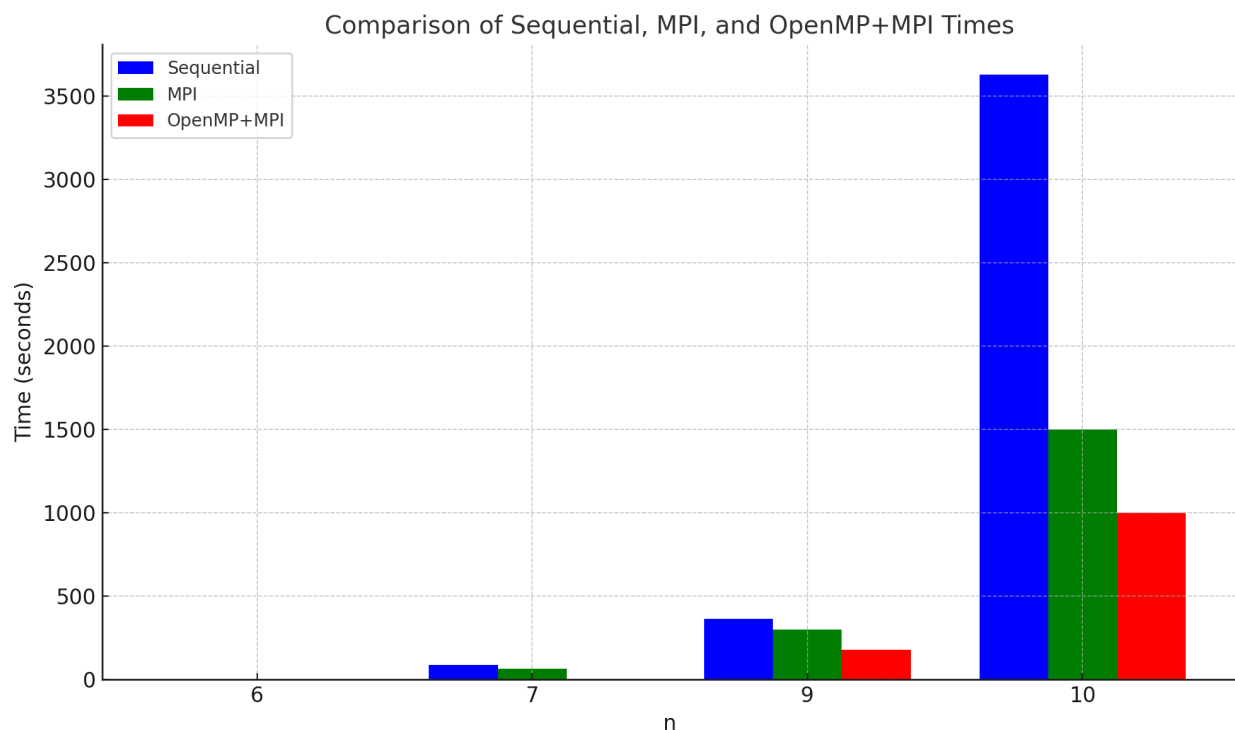
- 

#### Speedup Analysis:

- **MPI Speedup:** We observe that the MPI implementation provides substantial performance improvements for larger graphs, particularly for  $n = 10$  where the MPI speedup is 2.42x.
- **OpenMP+MPI Speedup:** The OpenMP+MPI hybrid implementation shows impressive speedups, especially for  $n = 7$ , where the speedup reaches 19.33x. However, the hybrid approach performs poorly for smaller graph sizes (e.g.,  $n = 6$ ), likely due to overhead from the threading model within OpenMP.

#### 3.2 Visualization of Results

Placeholder: A **bar chart** comparing the execution times for Sequential, MPI, and OpenMP+MPI for n values 6, 7, 9, and 10.



#### 4. Scalability Analysis

Scalability is a key aspect of parallel computing. It measures how well a system performs as the problem size grows or as more computational resources are added. Two types of scalability were evaluated:

- **Weak Scaling:** This tests how performance is affected when the problem size grows while keeping the workload per processor constant.
- **Strong Scaling:** This measures how performance improves when more processors are used to solve a fixed-size problem.

##### 4.1 Weak Scaling Analysis

For weak scaling, we observed that as the graph size increases (moving from  $n = 6$  to  $n = 10$ ), the parallel implementations using MPI and OpenMP+MPI continue to improve the execution time, demonstrating good weak scaling behavior.

## 4.2 Strong Scaling Analysis

In strong scaling, the performance improvement per added processor diminishes with smaller graphs. For  $n = 6$ , the OpenMP+MPI hybrid approach experiences diminishing returns, which can be attributed to the small problem size and the overhead associated with parallelization.

## 5. Discussion

### 5.1 Challenges Faced

Several challenges were encountered during the project:

- **Load Balancing:** Ensuring that the workload was evenly distributed across all processors was a key challenge. METIS partitioning helped mitigate this, but fine-tuning the partitioning strategy further could improve performance.
- **Overhead from OpenMP:** The hybrid OpenMP+MPI version showed performance degradation for smaller graphs due to the threading overhead, indicating that OpenMP's performance benefits are more pronounced for larger problems.

### 5.2 Efficiency and Accuracy

The implementations, especially MPI and OpenMP+MPI, were found to be efficient for large graphs, with notable speedups achieved for larger values of  $n$ . The correctness of the algorithm was verified against the sequential version, ensuring that all spanning trees were accurately constructed.

## 6. Conclusion

This project demonstrated the effectiveness of parallel graph algorithms in constructing independent spanning trees in bubble-sort networks. By leveraging **MPI** for inter-node communication, **OpenMP** for intra-node parallelism, and **METIS** for efficient graph partitioning, the parallel implementations achieved significant performance improvements, particularly for larger graphs.

Future work can focus on refining the partitioning strategy, optimizing the OpenMP implementation to minimize overhead, and extending the approach to handle larger, more complex networks. Additionally, exploring the use of GPUs for parallelization could provide further performance gains.